

Information Systems Institute

Distributed Systems Group (DSG)
VU Distributed Systems Technologies SS 2013 (184.260)

Assignment 1

Submission Deadline: 12.4.2013, 18:00

General Remarks

- Group work is not allowed in the lab. You have to work alone. Discussions with colleagues (e.g., in the TUWEL forum) are allowed, but the code has to be written alone. If we find that two students submitted the same, or very similar assignments, these students will be graded with 0 points (no questions asked). We will use automated plagiarism checks to compare solutions. Note that we will also include the submissions of previous years in our plagiarism checks.
- No deadline extensions are given. Start early and, after finishing your assignment, upload your submission as a zip file to TUWEL. If you think that you will be hard-pressed making the deadline you should upload a first version well before time runs out! We will grade whatever is there at the deadline, there is no possibility to submit later on.
- Make sure that your solution compiles and runs without errors. If you are unsure, test compiling your submission on different computers (e.g., one of the ZID lab computers). Preferred target platform for the entire lab (all 3 assignments) is JDK 7¹. If you decide to develop your code using JDK 6, some tweaks (endorsed libraries) may be required to get the project running under this version (especially later for assignments 2 and 3). Before you submit, make sure to test your solution with JDK 7 - points will be deducted if your code does not build/run out of the box.
- The assignment project is set up using Apache Maven2². We provide a project template that contains some code interfaces and JUnit tests which will assist you in developing your code. Please stick exactly to the provided interfaces as we will check your solutions in an automated test environment. The Maven project is split up into multiple modules which correspond to the different sub-parts of the assignments. **Note:** If all unit tests are passing in your solution, it does *not* necessarily mean that you will receive all the points. We will perform additional code checks and run tests that are not provided in the template. The tests included in the template should merely help you get started and assist you in developing your solution. If you want to further increase your test coverage, you may also add new unit tests, but this is optional. Do not modify any of the pre-defined tests - in any case, we will check your code with the original tests from the template.
- A good introduction into JPA can be found in Part VI of the Java EE 6 Tutorial³.
- Before asking questions about Hibernate, check out the Hibernate Documentation⁴ page. If you cannot find an answer to your question, consult the Hibernate forum⁵.
- Use MySQL 5.1+⁶ as your database management system.
- For the MongoDB tasks, consult the MongoDB Java tutorial⁷.
- We expect (as can be seen in the persistence-unit configuration of persistence.xml) that you setup a database `dst` that can be accessed by user `root` (without a password). You may of course change the settings of the configuration for your work at home, but please reset them to the original values in your submitted solution (and make sure it still works). Again: make sure that all settings are as expected before you submit!
- Please make sure to add reasonable logging output to help us keep track of what your solution does. No debug output is very bad, and too much (e.g., many screen pages) is just as bad. Aim for a good middle ground, which allows us to check your solution quickly.
- **Later assignments will be based on (parts of) the solution of this lab. Keep this in mind while implementing your code.**

¹<http://www.oracle.com/technetwork/java/javase/downloads/>

²<http://maven.apache.org/>

³<http://download.oracle.com/javaee/6/tutorial/doc/bnbp.html>

⁴<http://hibernate.org/docs>

⁵<https://forum.hibernate.org/>

⁶<http://dev.mysql.com/downloads/mysql/5.1.htm#downloads>

⁷<http://www.mongodb.org/display/DOCS/Java+Tutorial>

A. Code Part

Command for building/testing parts 1.-4.:	mvn install -Pass1-jpa
Command for building/testing part 5.:	mvn install -Pass1-nosql

1. Mapping Persistent Classes and Associations (16 Points)

Throughout this lab, it is your task to build the enterprise application for managing high-performance computing grids. A user can create and assign jobs to a grid. The grid then executes the job on a cluster of computers, i.e., on many machines simultaneously. Our simple assumption here is that each job needs a fixed number of CPUs to execute. Jobs may be assigned freely to computers and their CPUs. For instance, we assume that a job which needs to run on 8 CPUs can either run on 2 quad-cores or on 8 single core computer. Computers are organized in clusters. Every computer belongs to exactly one cluster. It is also possible that a cluster is recursively composed of other clusters. Each cluster is maintained by exactly one administrator.

Please note that your task in this and the following assignments is to build the Grid Management System, not the Grid itself. We will only simulate the actual execution of jobs.

The domain model that we use as starting point is defined in the following UML class diagram. We will come back and slightly extend this data model in future tasks and assignments.

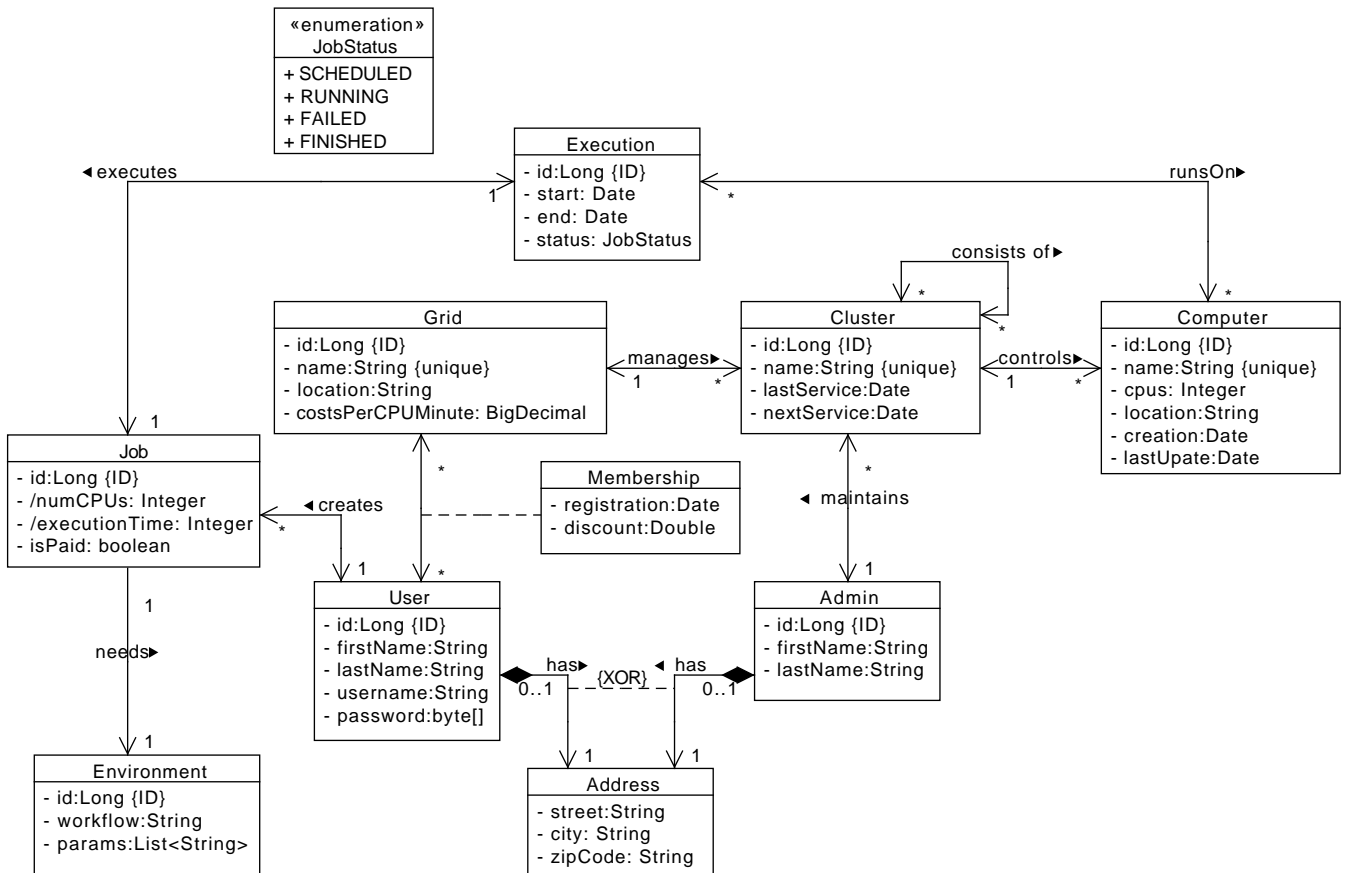


Figure 1: Grid - Domain Model

Note that each user and each admin has an address. These addresses shall be mapped as embedded classes. Also notice that each job needs some kind of environment. The environment consists of a workflow and a list of parameters. Obviously, the ordering of the parameters in this list is important, i.e., the order should not change after persisting and loading the entity. Jobs are executed on a number of computers in parallel (a job runs on many computers, but at all at the same time). Users do not need to be members of grids to assign jobs. Memberships are identified via the grid's and the user's ID. If the user has a membership, the grid also might give discounts (in percent) to users for assigned jobs. Finally, the password of users is stored as a cryptographic hash of a password string (you should never store plain text passwords to the database). For simplicity and due to its wide availability you should use MD5 as hashing function (although in practice MD5 is known to have some vulnerabilities). MD5 hashes are fixed size, so make sure that Hibernate uses the optimal data type for your password column.

In the template, **interfaces** (starting with capital letter "I") are provided for all **model classes** and **data access objects** (DAOs). You should create an implementation class for each of the interfaces and make sure to correctly instantiate objects using the **ModelFactory** and the **DAOFactory**. Put all the implementation classes you add into a separate **impl** sub-package, i.e., **dst.ass1.jpa.model.impl** and **dst.ass1.jpa.dao.impl**.

1.a. Basic Mapping

Map these classes and all associations using the Hibernate JPA implementation (i.e., only make use of the package `javax.persistence.*` in your mappings and the respective test code). For all classes, use annotations for the mappings. The only entity you should use XML mappings for is the **Computer** entity. Make sure that you implement navigation in the data model as specified in the class diagram (for instance, the job-execution association is bidirectional, so implement it that way).

Note that you will also need your solution for this part for almost all tasks in all later assignments, so we strongly recommend that you solve at least this part of the assignment.

1.b. Inheritance

While solving task (1.a.), you may have noticed that users and admins resemble each other in many ways. Therefore, we will now implement an inheritance relationship via the abstract entity *Person*. Expand and adapt your code from task (1.a.) by choosing one of the three well-known inheritance patterns. Be prepared to argue your choice during the practice lesson! At this point, we also add information about the user bank account to be able to debit money for the processed jobs. The required changes are shown in the following diagram. You also need to implement the two new unique constraints, as well as the new not-null constraint.

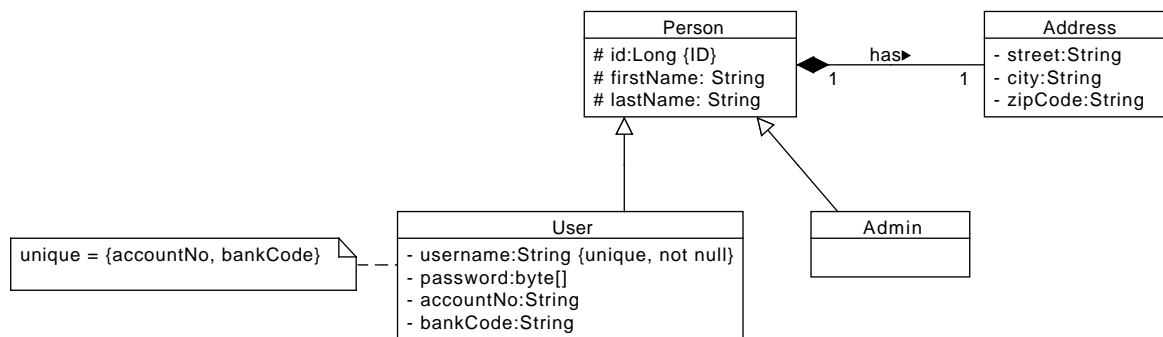


Figure 2: Grid - Inheritance Model

Now that you have finished mapping the grid's domain, you need to test/demonstrate your mapping.

The unit tests in the template store a reasonable amount of entities in the database, retrieve them, update some of them, and delete some of them again. However, you may extend the pre-defined tests and also include additional test classes to cover any special situations and corner cases, if necessary. Make sure that your solution works in its entirety, that is, no unnecessary tables get created, no unintended information is lost when deleting entities (e.g., because of unexpected cascades), and, of course, that there are no exceptions thrown when doing any of it.

2. JPA-QL and Hibernate Criteria API (8 Points)

2.a. Write the following **Named Queries** (no plain SQL, and no inline queries!) for the domain model implemented in Task 1. Your queries should solely accept the specified parameters and return the requested types. Make sure to name your queries exactly as specified in the pre-defined tests of the template.

- Find all users who have an active membership for a specific grid and assigned at least x jobs to it (given the name of the grid and the number of jobs).
- Find the most active user(s), i.e., the user(s) who assigned the largest amount of jobs.

2.b. Implement a named query to find all computers in Vienna (location field starts with 'AUT-VIE') and, for each of them, calculate the total usage (sum of end date minus start date for all executions of this computer). You do not need to find a single query to solve this task (you can use a combination of Java code and named queries), but you have to keep ORM-performance in mind, i.e., make sure that your solution is also reasonably fast if you have many executions on many computers. During the presentation you should be able to answer which problems you can run across here. For the unit tests to succeed, implement the corresponding DAO interfaces (methods are already defined in the interfaces).

2.c. Write the following queries using the **Hibernate Criteria API** (i.e., you are not able to conform to JPA this time). All parameters should be optional (i.e., indicated by `null` values). Again, put the code into your implementation of the corresponding DAO methods.

- Find all jobs that were created by a specified user which execute a given workflow (user is given by his username, and the workflow as workflow name).
- Use **Query-by-Example** to find all jobs with JobStatus `FINISHED` and a specified start and finish date.

2.d. Now we want to optimize the following query by adding an **index** with Hibernate:

- `"Select u from User u where u.password = MD5(?)"`

Assume that the user base will be very large, and this query will be executed often (i.e., whenever a user logs in). Assign the correct index value to speed up this query. Prepare to argue during the practice lessons, in which cases this index is helpful, and when it might even be harmful to the performance of the application. We encourage you to also run some large-scale tests (adding huge amounts of users, running the query, taking the time, deleting the users again) to check whether (and under which circumstances) the index has an influence on the performance. Implementing large-scale tests is optional, but could nicely contribute to our discussion during the lab session.

3. Bean Validation (3 Points)

The Bean Validation API is a feature in Java EE version 6 and above. You can use it to introduce restrictions on class fields or complete class instances, e.g., to verify that a certain String contains a well-formed email-address or that an Integer field has a value in the range between two specified constants. Bean Validation becomes especially useful when validating user input before storing it to the database, but it is explicitly not tied to any programming model or application layer.

3.a. In a first step we will start by applying some built-in constraints of the **Bean Validation API** to our computer (this time, you are allowed to use annotations for the computer entity):

- A computer has dates for creation and lastUpdate, both have to be in the past.
- Every computer has a name not shorter than 5 and no longer than 25 characters.
- The location of the computer starts with three upper-case alphabetic characters (for the Country) followed by a '-', another 3 upper-case alphabetic characters (for the City) followed by a '@' and at least 4 digits (for the zip-code). Example: 'AUT-VIE@1040'.

3.b. The grid only uses Computers which have a specific amount of CPUs. Therefore, provide a constraint annotation and a constraint validator to ensure that the amount of CPUs is between a defined upper and lower bound. If validation fails, return meaningful error messages. Put the code into **dst.ass1.jpa.validator**. The annotation should look like this:

```
@CPUs(min = 4, max = 8)
```

The pre-defined tests in the template instantiate valid and invalid computer entities, and use an instance of `javax.validation.Validator` to validate each entity. Invalid entities should violate every single constraint. We advise you to also create your own test class with additional test cases to verify each constraint individually.

4. Entity Manager, Entity Listeners and Interceptor (5 Points)

4.a. Entity Manager - Lifecycle

Up to this point, you already gained some experience storing and retrieving entities using the EntityManager. Now it is time to study the lifecycle of Persistent Entities. Create a transaction sequence to illustrate the persistence lifecycle of a job object (there has to be an association with a user). Additionally, make some source code comments that describe the current state of the activity object. Put the resulting code into **dst.ass1.jpa.lifecycle.EMLifecycleDemo**. Make sure that all possible states are covered by your example.

4.b. Entity Listener

Entity Listeners allow us to register callback methods for specified lifecycle events. Provide an entity listener which sets the creation- and lastUpdate-field every time a new computer gets stored and overwrites the lastUpdate-field every time an already existing computer entity is updated. Implement your callback methods in the class **dst.ass1.jpa.listener.ComputerListener**.

4.c. Default Listener

Now your task is to implement a simple default listener. This listener should record the number of entities loaded, updated and removed. This listener also records how often a persist operation was successfully called and how much time all store-operations took in total. Implement your listener by completing the class **dst.ass1.jpa.listener.DefaultListener**. Make sure that your listener implementation is thread-safe!

4.d. Interceptor

Task 2.b. aimed at minimizing the number of selects issued to the database. Now we will use the Hibernate Interceptor interface to count the number of database select statements that Hibernate actually uses in the background. The Hibernate Interceptor interface provides callbacks from the session to the application, allowing the application to inspect and/or manipulate properties for a specific event. Implement the provided **dst.ass1.jpa.interceptor.SQLInterceptor** and count the number of selects for computers and executions. It is ok to check if the SQL-String starts with the respective select-statement. In addition, you should provide methods to reset the select statement counter. Use the interceptor to validate that your solution for Task 2.b. works efficiently.

5. NoSQL Database (8 Points)

So far, all tasks made use of a traditional SQL database (MySQL) and JPA/Hibernate. For the last practical tasks of this assignment, we will look at a different database paradigm, so-called NoSQL (Not only SQL) databases, exemplified on the basis of *MongoDB*. The project template already contains a dependency to MongoDB, and upon initialization the class **TestData** runs an embedded MongoDB server. That is, no local installation is required on your machine. The template project uses the default MongoDB port 27017 (make sure this port is free), and skips any authentication features (evidently, you would not want to do this in production, but for our tests this is the easiest way to get started).

The package **dst.ass1.nosql** in the project template contains an interface for loading test data into MongoDB (**IMongoDbDataLoader**), and one interface for running queries against the MongoDB database (**IMongoDbQuery**). Use these interfaces for the remaining parts of this section, and put your implementation classes into a sub-package **dst.ass1.nosql.impl**. Also make sure to properly instantiate and return your implementation classes in **MongoDbFactory**.

5.a Storing workflow results in MongoDB

So far, all data that we wanted to store in our Grid management system was fairly well-structured and uniform. However, assume now that we also need to capture and store the output produced by Grid workflows. These data are evidently quite unstructured, as different types of workflows will produce different types of output. For instance, a bioinformatics workflow might produce some representation of chromosome data, while a mathematical workflow might deliver some matrices as result. Furthermore, there is no way at design time to know which data we will need to store. Hence, we will use MongoDB, a schema-free NoSQL database for this part of the system.

Every workflow output can be represented as a JSON document with 2 or more properties. The first property is **job_id**, which refers back to the relational database ID of the job that produced this output. The second property is **last_updated**, containing a UNIX timestamp when this document was last changed. Finally, workflow outputs can have 0 or more additional properties, containing the actual payload of the document. An example is provided below.

```
{
  "job_id" : 2 ,
  "last_updated" : 1329469221971,
  "result_matrix" : {
    "matrix" : [ [ 1,0,0,0] , [0,1,0,0] , [0,0,1,0] , [0,0,0,1]]
  } ,
  "type" : "identity_matrix"
}
```

Your first task is to add output documents for every finished job to your MongoDB installation. Add at least 5 different, not entirely trivial and at least partly structurally different, documents. To get you started, the provided template contains 3 example documents in JSON notation (see class **MongoTestData**).

Put the code into your implementation of the method **IMongoDbDataLoader.loadData**. Inside the **loadData** method, you should first load all finished jobs from the MySQL database (using the named query created earlier in this assignment), loop over all jobs and store a JSON object with the job output to MongoDB. Use the official MongoDB Java driver (already included in the Maven dependencies of the template). Since it does not really matter which output you assign to which job, you can select any random result data from the JSON documents in **MongoTestData**. Make sure to correctly set the **job_id** and **last_updated** attributes in all JSONs you save to MongoDB. Furthermore, as we will often retrieve documents by job id, you should add an index to speed up such queries.

5.b Querying workflow results

MongoDB primarily uses the query-by-example pattern for retrieving data, which we have already used in Task 2.c. In your implementation of the interface **IMongoDbQuery**, you have to write

two simple queries, which (1) fetch the result document for a given job id (e.g., for the job id 1) and (2) fetch all results with a last updated field later than 1325397600 (Unix timestamp of 01.01.2012). For the second query, use a filter to retrieve only the actual payload of the document (no other properties should be contained in the delivered result). Provide reasonable logging output with information about the results of your queries.

5.c Map/Reduce queries

In addition to query-by-example, we can also use the more powerful, but also more complex, map/reduce pattern for data retrieval and processing. MongoDB supports map/reduce via JavaScript, that is, you will need to write both the map and the reduce functions as JavaScript functions that operate on your stored documents. You can find a good example online⁸. Another blog post shows how you can call map/reduce via the Java driver⁹.

Your task is now to implement the method **IMongoDbQuery.mapReduceWorkflow** to use map/reduce for reporting some rudimentary statistics about the stored workflow output documents. The output of your map/reduce query should be a list of all existing (top-level) payload document properties, along with the number of documents that they occur in. Do not report the properties containing the job id or the last updated field. Similarly, skip the `_id` field which MongoDB automatically generates for each document. Essentially, this report allows us to reason about how many documents there are for each type of workflow. See below for an example output. As before, print the results of your query to the system output stream.

```
{ "_id" : "alignment_block" , "value" : 3}
{ "_id" : "logs" , "value" : 1}
{ "_id" : "result_matrix" , "value" : 4}
{ "_id" : "type" , "value" : 4}
```

⁸<http://blog.fiesta.cc/post/10980328832/walkthrough-a-mongodb-map-reduce-job>

⁹http://blog.evilmongrelabs.com/2011/02/28/MongoDB-1_8-MR-Java/

B. Theory Part

The following questions will be discussed during the practice lesson. At the beginning of the each lesson we hand out a list where you can specify which questions you have prepared and are willing to present. We will then select students at random who checked a question to discuss the question (you know the procedure from your math courses). If you are asked to discuss a question but fail to provide a strong answer, you will lose **all** points for the theory part of this assignment.

5. Annotations vs. XML (1 Point)

In the previous tasks you already gained some experiences using annotations and XML. What are the pros and cons of each approach, when would you use which one? Answering these questions, also keep in mind maintainability and the different roles usually involved in software projects.

6. Versioning (1 Point)

JPA provides a feature called versioning. Why and under which circumstances can this feature be useful? Think about a situation where optimistic locking may result in an (desired) exception.

7. Read-Locks (2 Point)

The EntityManager allows the programmer to set Read-locks on specified objects. What are the consequences on concurrent threads when one thread sets such a lock? Think about use-cases this behaviour may be adequate for. What problems can arise?

8. Database Isolation-Levels (2 Points)

Have a look at the different isolation-levels modern databases provide. What kind of problems might occur due to concurrency issues at what kind of level? Is it really necessary to protect your application against every type of flaw?

9. Database Indices (2 Point)

What is the purpose and functioning of a database index? Using which data structures do database management systems typically store an index internally, and what are important characteristics of these data structures? What is the basic tradeoff of using an index, what are its limitations? Think of two concrete examples - one in which an index leads to an improvement and one in which the index is useless (i.e., does not lead to an improvement).

10. NoSQL Databases (2 Point)

What general types of NoSQL databases exist? Name prominent examples for each type of database, and argue when you should be using them (and also, when you should specifically *not* use them).